

Project Atlas – Q4 Executive Status

Building the next-generation AI platform for a more productive enterprise



ON TRACK

Overall project health
Unchanged since Q3

Completion
72%
of scope delivered




Budget used
64%
\$4.1M of \$6.4M




BOTTOM LINE
MVP ships on time. GPU capacity is the one risk to a Q2 launch.

Workstream progress

% complete · status



Key milestones



KEY RISK Medium

GPU capacity from our third-party provider is not yet confirmed for Q1

If unresolved, the public beta slips four to six weeks



NEXT DECISION By Dec 6

Approve a \$0.6M reserve for dedicated GPU capacity

Secures the Feb 15 beta and de-risks the Q2 launch

Three workstreams. One critical dependency.

Product is ready; engineering and go-to-market converge on the beta gate.

Product

82%



Validated scope and UX

Owner: Maya Chen

Engineering

74%



API, inference and hardening

Owner: Leo Martins

Go-to-Market

58%



Enablement and beta demand

Owner: Priya Shah

Shared dependency: GPU capacity must be secured before integrated beta testing.

Delivery is ahead of budget consumption.

A controlled reserve can protect the beta without changing the approved envelope.

\$6.4M

Approved envelope

Total project funding

\$4.1M

Used to date


64% of approved budget

\$2.3M

Remaining

Includes a proposed \$0.6M reserve

Scope delivered



72%

Budget consumed



64%

A narrow risk window. A practical mitigation.

Treat capacity as a sourcing decision with a clear owner and deadline.

MEDIUM RISK

GPU allocation remains unconfirmed.

Exposure: public beta slips four to six weeks if supply is not secured.

Owner: Engineering + Procurement

DEC 2

Validate capacity

Confirm regional supply and performance.

DEC 6

Approve reserve

Commit \$0.6M for dedicated allocation.

DEC 12

Close contract

Secure supply before the beta test gate.

Approve the reserve. Protect the launch.

Decision required by December 6, 2026. Proposed release remains Q2 2027.

\$0.6M

Capacity reserve

Dedicated GPUs for beta and scale-up

Approval conditions

Capacity confirmed in writing

Contract stays within approved budget

Performance validated before release

Next gate: December 15 MVP release, followed by the February 15 public beta.